

Project "Boston housing"

Stanin Vladislav

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm

from scipy.stats import ttest_ind, ttest_rel, mannwhitneyu, pearsonr
```

Data installation

```
In [2]: data = pd.read_csv('../data/BostonHousing.csv')
data
```

Out[2]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
...
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	9.67	22.4
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08	20.6
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64	23.9
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48	22.0
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90	7.88	11.9

506 rows × 14 columns

little EDA

```
In [3]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   crim        506 non-null    float64
1   zn          506 non-null    float64
2   indus       506 non-null    float64
3   chas        506 non-null    int64  
4   nox         506 non-null    float64
5   rm          506 non-null    float64
6   age         506 non-null    float64
7   dis         506 non-null    float64
8   rad         506 non-null    int64  
9   tax         506 non-null    int64  
10  ptratio     506 non-null    float64
11  b           506 non-null    float64
12  lstat       506 non-null    float64
13  medv        506 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB
```

```
In [4]: data.isna().sum()
```


Out[4]:

crim	0
zn	0
indus	0
chas	0
nox	0
rm	0
age	0
dis	0
rad	0
tax	0
ptratio	0
b	0
lstat	0
medv	0
dtype:	int64

No problems with data

Distribution of target variable (*medv*):

```
In [5]: sns.histplot(data[['medv']], stat='density', color='blue', legend=False);  
plt.xlabel('medv');
```



Looks like normal

```
In [6]: predictors = data.iloc[:, :-1]  
medv = data.iloc[:, -1]
```

Standartization

```
In [7]: means = predictors.mean(axis=0)  
stds = predictors.std(axis=0)  
  
scaled_preds = (predictors - means) / stds  
  
scaled_preds.head()
```

```
Out [7]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	
0	-0.419367	0.284548	-1.286636	-0.272329	-0.144075	0.413263	-0.119895	0.140075	-0.981871	-0.665949	-1.457558	0.440616	-1.074
1	-0.416927	-0.487240	-0.592794	-0.272329	-0.739530	0.194082	0.366803	0.556609	-0.867024	-0.986353	-0.302794	0.440616	-0.49
2	-0.416929	-0.487240	-0.592794	-0.272329	-0.739530	1.281446	-0.265549	0.556609	-0.867024	-0.986353	-0.302794	0.396035	-1.20
3	-0.416338	-0.487240	-1.305586	-0.272329	-0.834458	1.015298	-0.809088	1.076671	-0.752178	-1.105022	0.112920	0.415751	-1.36
4	-0.412074	-0.487240	-1.305586	-0.272329	-0.834458	1.227362	-0.510674	1.076671	-0.752178	-1.105022	0.112920	0.440616	-1.02

First linear model

```
In [8]: X = sm.add_constant(scaled_preds)  
model_scaled = sm.OLS(medv, X)  
results_scaled = model_scaled.fit()  
  
print(results_scaled.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	medv	R-squared:	0.741			
Model:	OLS	Adj. R-squared:	0.734			
Method:	Least Squares	F-statistic:	108.1			
Date:	Tue, 13 Dec 2022	Prob (F-statistic):	6.72e-135			
Time:	14:51:45	Log-Likelihood:	-1498.8			
No. Observations:	506	AIC:	3026.			
Df Residuals:	492	BIC:	3085.			
Df Model:	13					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	22.5328	0.211	106.814	0.000	22.118	22.947
crim	-0.9291	0.283	-3.287	0.001	-1.484	-0.374
zn	1.0826	0.320	3.382	0.001	0.454	1.712
indus	0.1410	0.422	0.334	0.738	-0.688	0.970
chas	0.6824	0.219	3.118	0.002	0.252	1.112
nox	-2.0588	0.443	-4.651	0.000	-2.928	-1.189
rm	2.6769	0.294	9.116	0.000	2.100	3.254
age	0.0195	0.372	0.052	0.958	-0.711	0.750
dis	-3.1071	0.420	-7.398	0.000	-3.932	-2.282
rad	2.6649	0.578	4.613	0.000	1.530	3.800
tax	-2.0788	0.634	-3.280	0.001	-3.324	-0.834
ptratio	-2.0626	0.283	-7.283	0.000	-2.619	-1.506
b	0.8501	0.245	3.467	0.001	0.368	1.332
lstat	-3.7473	0.362	-10.347	0.000	-4.459	-3.036
=====						
Omnibus:	178.041	Durbin-Watson:	1.078			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	783.126			
Skew:	1.521	Prob(JB):	8.84e-171			
Kurtosis:	8.281	Cond. No.	9.82			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The *age* and *indus* predictors not significantly predicts the *medv*

Checking model without these predictors:

```
In [9]: model_new = sm.OLS(medv, X.drop(columns=["indus", "age"]))
results_new = model_new.fit()

print(results_new.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	medv	R-squared:	0.741			
Model:	OLS	Adj. R-squared:	0.735			
Method:	Least Squares	F-statistic:	128.2			
Date:	Tue, 13 Dec 2022	Prob (F-statistic):	5.54e-137			
Time:	14:51:45	Log-Likelihood:	-1498.9			
No. Observations:	506	AIC:	3022.			
Df Residuals:	494	BIC:	3072.			
Df Model:	11					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	22.5328	0.211	107.018	0.000	22.119	22.946
crim	-0.9325	0.282	-3.307	0.001	-1.486	-0.379
zn	1.0692	0.315	3.390	0.001	0.450	1.689
chas	0.6905	0.217	3.183	0.002	0.264	1.117
nox	-2.0135	0.410	-4.915	0.000	-2.818	-1.209
rm	2.6711	0.285	9.356	0.000	2.110	3.232
dis	-3.1432	0.391	-8.037	0.000	-3.912	-2.375
rad	2.6088	0.552	4.726	0.000	1.524	3.693
tax	-1.9850	0.568	-3.493	0.001	-3.102	-0.868
ptratio	-2.0492	0.279	-7.334	0.000	-2.598	-1.500
b	0.8482	0.244	3.475	0.001	0.369	1.328
lstat	-3.7316	0.339	-11.019	0.000	-4.397	-3.066
=====						
Omnibus:	178.430	Durbin-Watson:	1.078			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	787.785			
Skew:	1.523	Prob(JB):	8.60e-172			
Kurtosis:	8.300	Cond. No.	7.90			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Model did not changed significantly arfter removing those predictors.

Nevertheless, F-statistics and R-adj were increased (and p-val decreased).


In the rest of report i will name these paramters as **Statistics** (If they become better, I will say that Statistics *increased*)

```
In [10]: X_new = X.drop(columns=["indus", "age"])
```

Checking model

Linear relationship


```
In [11]: fig = sns.pairplot(pd.concat([scaled_preds.drop(columns=["indus", "age"]), medv], axis=1), kind="reg", corner=True);  
fig.savefig("pairplot.png");
```



Most of predictors have linear relationship with *medv*, but *lstat* which have looking like exponential relationship

Let's check:

```
In [12]: sns.scatterplot(y=predictors.lstat,  
                        x=medv);
```




```
In [13]: print(f"medv and lstat correlated with Pearson's r value {pearsonr(y=predictors.lstat, x=medv).statistic} and pvalue {pearsonr(y=predictors.lstat, x=medv).pvalue}")
```

medv and lstat correlated with Pearson's r value -0.737662726174015 and pvalue $5.081103394386392e-88$

Checking relationship between *lstat* and $-\ln(\text{medv})$

```
In [14]: sns.scatterplot(y=predictors.lstat,
                        x=-np.log(medv+1));
```



```
In [15]: print(f'ln(medv) and lstat correlated with Pearson's r value {pearsonr(y=predictors.lstat, x=np.log(medv+1)).statistic}')
ln(medv) and lstat correlated with Pearson's r value  $-0.8043$  and pvalue  $5.230310856128727e-116$ 
```

Since Pearson's r not changed significantly, we can leave it

However,

Additional check of model without *lstat* predictor:

```
In [16]: model_wo_lstat = sm.OLS(medv, X_new.drop(columns=["lstat"]))
results_wo_lstat = model_wo_lstat.fit()

print(results_wo_lstat.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	medv		R-squared:	0.677		
Model:	OLS		Adj. R-squared:	0.670		
Method:	Least Squares		F-statistic:	103.7		
Date:	Tue, 13 Dec 2022		Prob (F-statistic):	1.33e-114		
Time:	14:52:35		Log-Likelihood:	-1554.5		
No. Observations:	506		AIC:	3131.		
Df Residuals:	495		BIC:	3177.		
Df Model:	10					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	22.5328	0.235	95.979	0.000	22.072	22.994
crim	-1.4336	0.310	-4.621	0.000	-2.043	-0.824
zn	1.0549	0.352	3.000	0.003	0.364	1.746
chas	0.7784	0.242	3.220	0.001	0.303	1.253
nox	-2.9458	0.447	-6.591	0.000	-3.824	-2.068
rm	4.2960	0.273	15.761	0.000	3.760	4.831
dis	-2.7631	0.434	-6.361	0.000	-3.617	-1.910
rad	2.6869	0.616	4.365	0.000	1.478	3.896
tax	-2.1965	0.633	-3.468	0.001	-3.441	-0.952
ptratio	-2.3219	0.310	-7.482	0.000	-2.932	-1.712
b	1.2392	0.269	4.602	0.000	0.710	1.768
=====						
Omnibus:	247.217		Durbin-Watson:	0.948		
Prob(Omnibus):	0.000		Jarque-Bera (JB):	2087.021		
Skew:	1.949		Prob(JB):	0.00		
Kurtosis:	12.154		Cond. No.	7.46		
=====						

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


Statistics of model decreased, therefore we leave *lstat* predictor in model

Checking if distribution of residuals is normal

```
In [17]: prediction = results_new.get_prediction(X_new)
medv_predicted = prediction.predicted_mean


In [18]: residuals = medv - medv_predicted

In [19]: sns.histplot(residuals, color='lightblue', binwidth=1);
plt.xlabel('Residuals');
```



Seems like normal...

```
In [20]: sm.qqplot(residuals, line='s');
```



Still looks like normal, with some deviations (richest houses)
but let's check these deviations:

Checking deviations

(Calculating Cook's distances)


```
In [21]: influence = results_new.get_influence()
cooks = influence.cooks_distance
n_deviations = (cooks[1] < 0.05).sum()
print(f'There are {n_deviations} significant deviations')
```

There are 0 significant deviations

There are no deviations!

Checking homoscedacity

```
In [22]: sns.scatterplot(x = medv, y = residuals);
plt.hlines(0,0, 51, color = 'red');
```



Only some very expensive houses deviate. Nevertheless, we can consider that the model has homoscedasticity

Checking VIF

```
In [23]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [24]: def vif(preds):
    vif_data = pd.DataFrame()
    vif_data["Predictor"] = preds.columns
    vif_data["VIF"] = [variance_inflation_factor(preds.values, i) for i in range(len(preds.columns))]
    return vif_data
```

```
In [25]: vif(X_new)
```

Out[25]:

	Predictor	VIF
0	const	1.000000
1	crim	1.789704
2	zn	2.239229
3	chas	1.059819
4	nox	3.778011
5	rm	1.834806
6	dis	3.443420
7	rad	6.861126
8	tax	7.272386
9	ptratio	1.757681
10	b	1.341559
11	lstat	2.581984

There are several predictors that is linked with others

```
In [26]: X_new_2 = X_new.drop(columns=["tax"])
vif(X_new_2)
```

Out[26]:

	Predictor	VIF
0	const	1.000000
1	crim	1.787963
2	zn	2.154054
3	chas	1.052428
4	nox	3.564036
5	rm	1.806735
6	dis	3.410587
7	rad	2.776775
8	ptratio	1.717222
9	b	1.338982
10	lstat	2.579040

It looks much better
VIF of *rad* reduced significantly and it means thar *tax* and *rad* was linked

Checking new model without *tax*

```
In [27]: model_new_2 = sm.OLS(medv, X_new_2)
results_new_2 = model_new_2.fit()

print(results_new_2.summary())
```


OLS Regression Results						
=====						
Dep. Variable:	medv		R-squared:	0.734		
Model:	OLS		Adj. R-squared:	0.729		
Method:	Least Squares		F-statistic:	136.7		
Date:	Tue, 13 Dec 2022		Prob (F-statistic):	1.84e-135		
Time:	14:52:36		Log-Likelihood:	-1505.0		
No. Observations:	506		AIC:	3032.		
Df Residuals:	495		BIC:	3079.		
Df Model:	10					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	22.5328	0.213	105.828	0.000	22.114	22.951
crim	-0.9018	0.285	-3.164	0.002	-1.462	-0.342
zn	0.8544	0.313	2.731	0.007	0.240	1.469
chas	0.7538	0.219	3.448	0.001	0.324	1.183
nox	-2.3540	0.402	-5.850	0.000	-3.145	-1.563
rm	2.7944	0.286	9.754	0.000	2.232	3.357
dis	-3.0098	0.394	-7.647	0.000	-3.783	-2.236
rad	1.1212	0.355	3.157	0.002	0.423	1.819
ptratio	-2.1972	0.279	-7.867	0.000	-2.746	-1.648
b	0.8856	0.247	3.591	0.000	0.401	1.370
lstat	-3.7715	0.342	-11.019	0.000	-4.444	-3.099
=====						
Omnibus:	166.907		Durbin-Watson:	1.090		
Prob(Omnibus):	0.000		Jarque-Bera (JB):	684.418		
Skew:	1.441		Prob(JB):	2.40e-149		
Kurtosis:	7.915		Cond. No.	5.02		
=====						


Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Statistics increased!
Now model looks pretty!

The largest modulo value of coefficient belongs to *lstat* predictor


rm vs raw *medv*

```
In [28]: sns.scatterplot(x = medv, y = predictors.lstat);
```



lstat vs predicted *medv*

```
In [29]: sns.scatterplot(x = medv_predicted, y = predictors.lstat);
plt.xlabel('medv_predicted');
```



OPTIONAL

Let's see parameters of suburbs that have *medv* variable > mean + 1 sd - It will be named *data_medv_high*

And suburbs that have *medv* variable < mean + 1 sd - It will be named *data_medv_low*

```
In [30]: mean_medv = medv.mean()
std_medv = medv.std()
```

```
In [31]: treschold = mean_medv + std_medv
```

```
In [32]: data_medv_high = data.loc[data.medv > treschold]
data_medv_low = data.loc[data.medv <= treschold]
```

Looking through correlation between medv and all predictors:

```
In [33]: (pd.concat([medv, X_new_2.drop(columns=['const'])], axis=1).corr().iloc[:,0])
```

```
Out[33]: medv      1.000000
crim      -0.388305
zn         0.360445
chas       0.175260
nox        -0.427321
rm         0.695360
dis        0.249929
rad        -0.381626
ptratio    -0.507787
b          0.333461
lstat      -0.737663
Name: medv, dtype: float64
```

I will check predictors with high correclatioon values

One-sided Mann-Whitney U test for all predictors vs *medv*:

```
In [34]: for col in data_medv_high.columns:
print(col, '- in cheap suburbs greater')
print(f"p-value - {mannwhitneyu(data_medv_low[col], data_medv_high[col], alternative='greater').pvalue}")
print(col, '- in expensive suburbs greater')
print(f"p-value - {mannwhitneyu(data_medv_low[col], data_medv_high[col], alternative='less').pvalue}")
print()
```

```
crim - in cheap suburbs greater
p-value - 6.192291924409378e-07
crim - in expensive suburbs greater
p-value - 0.9999993835304415

zn - in cheap suburbs greater
p-value - 0.9999999999961751
zn - in expensive suburbs greater
p-value - 3.8556018089623235e-12

indus - in cheap suburbs greater
p-value - 1.062521626251081e-14
indus - in expensive suburbs greater
p-value - 0.999999999999895

chas - in cheap suburbs greater
p-value - 0.9992560575911941
chas - in expensive suburbs greater
p-value - 0.000749130155310187

nox - in cheap suburbs greater
p-value - 3.640067266843077e-07
nox - in expensive suburbs greater
p-value - 0.9999996376482426

rm - in cheap suburbs greater
p-value - 1.0
rm - in expensive suburbs greater
p-value - 9.790460577535401e-33

age - in cheap suburbs greater
p-value - 0.00014914106723406078
age - in expensive suburbs greater
p-value - 0.9998513685389849

dis - in cheap suburbs greater
p-value - 0.9816522883378049
dis - in expensive suburbs greater
p-value - 0.01838761393683482

rad - in cheap suburbs greater
p-value - 0.0001496462931231421
rad - in expensive suburbs greater
p-value - 0.9998508754647684

tax - in cheap suburbs greater
p-value - 1.828698636468821e-10
tax - in expensive suburbs greater
p-value - 0.9999999998181772

ptratio - in cheap suburbs greater
p-value - 1.9715109522953002e-15
ptratio - in expensive suburbs greater
p-value - 0.999999999999998

b - in cheap suburbs greater
p-value - 0.9325626166258136
b - in expensive suburbs greater
p-value - 0.0675538449832815

lstat - in cheap suburbs greater
p-value - 1.1394403179327003e-31
lstat - in expensive suburbs greater
p-value - 1.0

medv - in cheap suburbs greater
p-value - 1.0
medv - in expensive suburbs greater
p-value - 5.358202079577759e-41
```

I'm very picky and suggest that predictors *b*, *rad*, *dis*, *age* and *chas* not significantly affects the cost of house

Testing *lstat*

Also, according to summary of last model, the most important parameter for higher cost is lower status of the population (percent)

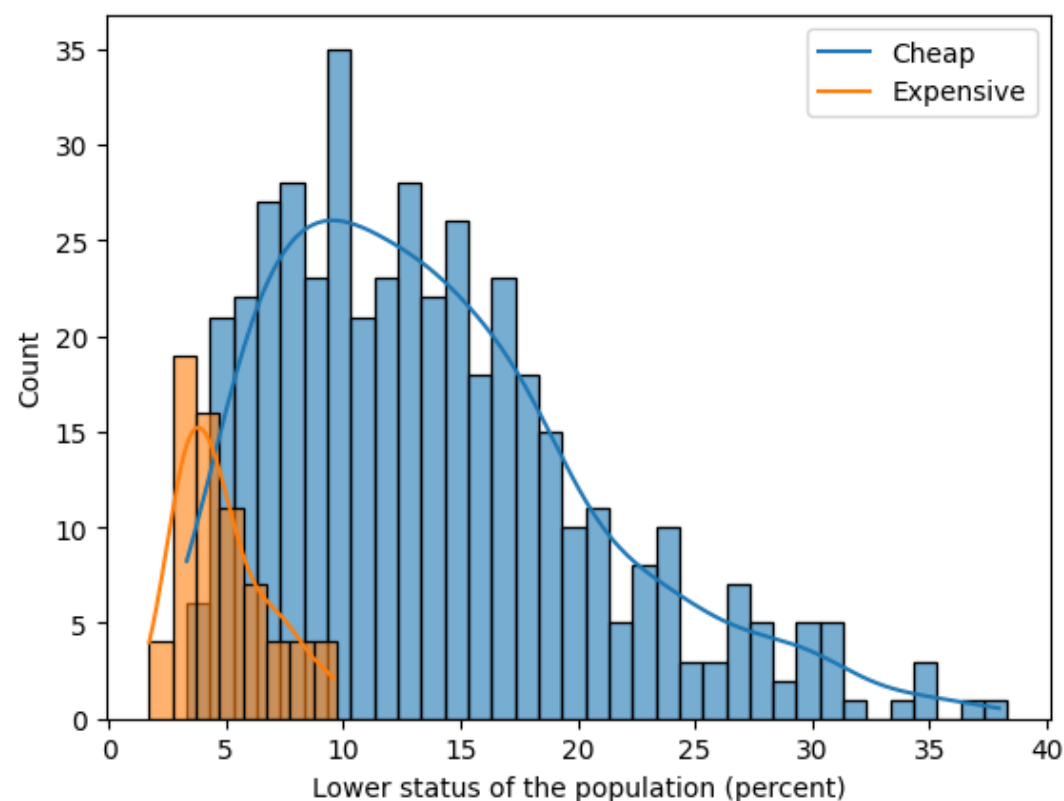
```
In [35]: data_medv_low.lstat.describe()

Out[35]: count    437.000000
         mean     13.883547
         std      6.880995
         min      3.330000
         25%      8.610000
         50%     12.800000
         75%     17.640000
         max     37.970000
         Name: lstat, dtype: float64
```

```
In [36]: data_medv_high.lstat.describe()
```

```
Out[36]: count    69.000000
mean      4.860000
std       1.941909
min       1.730000
25%       3.320000
50%       4.450000
75%       6.050000
max       9.590000
Name: lstat, dtype: float64
```

```
In [37]: b1 = sns.histplot(data_medv_low.lstat, kde=True,
                           alpha=0.6, binwidth=1);
b2 = sns.histplot(data_medv_high.lstat, kde=True,
                  alpha=0.6, binwidth=1);
plt.xlabel('Lower status of the population (percent)');
plt.legend(['Cheap', 'Expensive']);
```



```
In [38]: pval_mann_lstat = mannwhitneyu(data_medv_high.lstat, data_medv_low.lstat, alternative='less').pvalue
print('Expensive suburbs have greater value of lower status of the population (percent) then between cheap suburbs\n
      f'with pvalue of Mann-Whitney U test: {pval_mann_lstat}')
```

Expensive suburbs have greater value of lower status of the population (percent) then between cheap suburbs
with pvalue of Mann-Whitney U test: 1.1394403179327003e-31

Testing *rm*

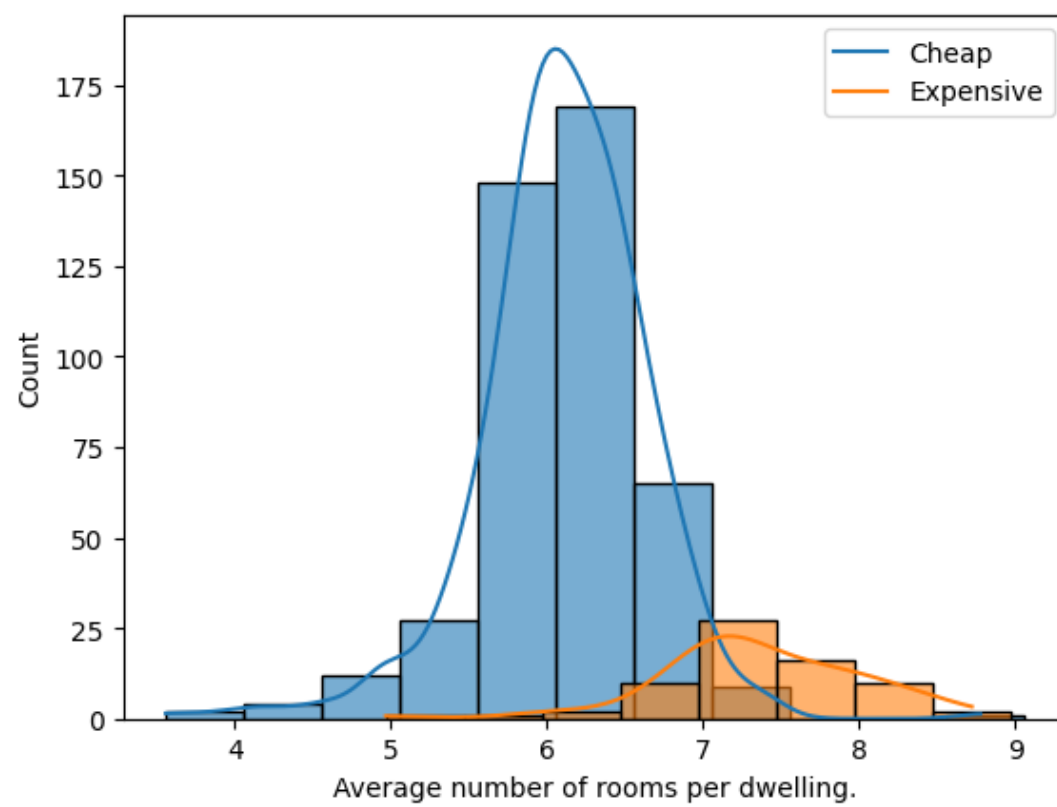
```
In [39]: data_medv_low.rm.describe()
```

```
Out[39]: count    437.000000
mean      6.112847
std       0.536631
min       3.561000
25%       5.857000
50%       6.127000
75%       6.431000
max       8.780000
Name: rm, dtype: float64
```

```
In [40]: data_medv_high.rm.describe()
```

```
Out[40]: count    69.000000
mean      7.372623
std       0.655013
min       4.970000
25%       6.998000
50%       7.267000
75%       7.820000
max       8.725000
Name: rm, dtype: float64
```

```
In [41]: b1 = sns.histplot(data_medv_low.rm, kde=True,
                           alpha=0.6, binwidth=0.5);
b2 = sns.histplot(data_medv_high.rm, kde=True,
                  alpha=0.6, binwidth=0.5);
plt.legend(['Cheap', 'Expensive']);
plt.xlabel('Average number of rooms per dwelling.');
```



```
In [42]: pval_mann_rm = mannwhitneyu(data_medv_low.dis, data_medv_high.rm, alternative='less').pvalue
print('Expensive suburbs have greater average number of rooms per dwelling than cheap ones',
      f'with one sided Mann-Whitney U test p-value: {pval_mann_rm}')
```

Expensive suburbs have greater average number of rooms per dwelling than cheap ones with one sided Mann-Whitney U test p-value: 7.653663514039866e-30

Testing *ptratio*

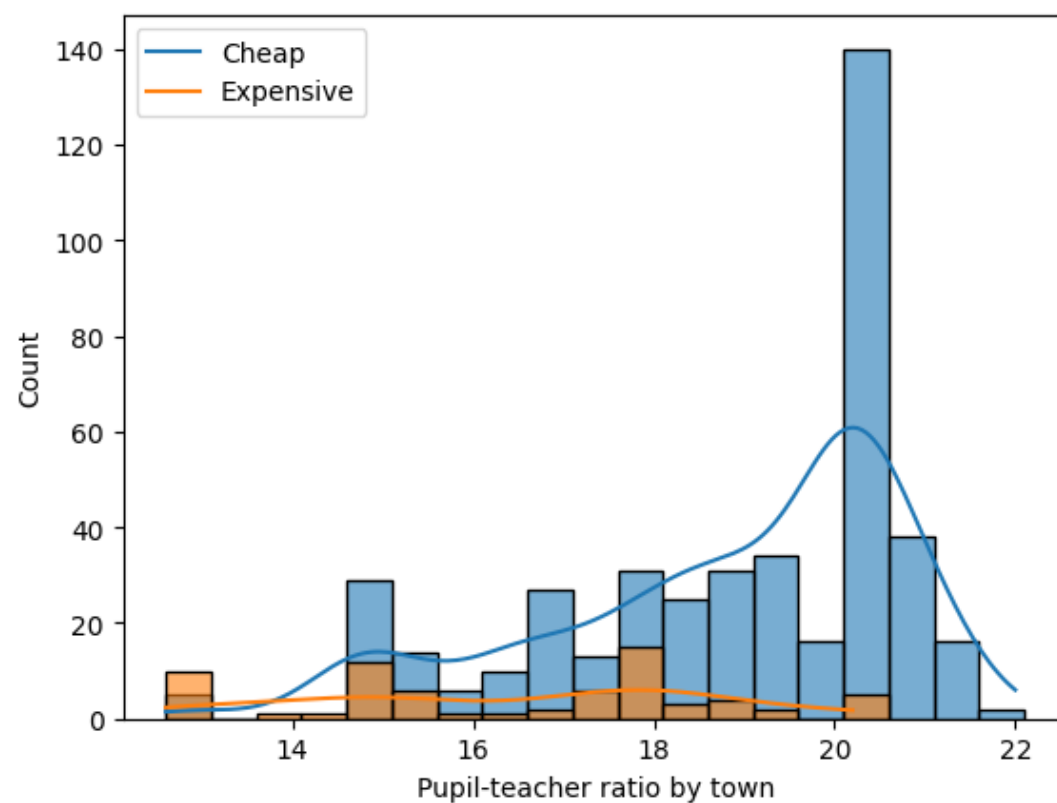
```
In [43]: data_medv_low.ptratio.describe()
```

```
Out[43]: count    437.000000
mean      18.778490
std       1.976883
min       12.600000
25%      17.800000
50%      19.200000
75%      20.200000
max       22.000000
Name: ptratio, dtype: float64
```

```
In [44]: data_medv_high.ptratio.describe()
```

```
Out[44]: count     69.000000
mean      16.410145
std       2.198806
min       12.600000
25%      14.700000
50%      17.400000
75%      17.900000
max       20.200000
Name: ptratio, dtype: float64
```

```
In [45]: sns.histplot(data_medv_low.ptratio, kde=True, alpha=0.6, binwidth=0.5);
sns.histplot(data_medv_high.ptratio, kde=True, alpha=0.6, binwidth=0.5);
plt.legend(['Cheap', 'Expensive']);
plt.xlabel('Pupil-teacher ratio by town');
# plt.ylim(0, 60);
```



```
In [46]: pval_mann_ptratio = mannwhitneyu(data_medv_low.pratio, data_medv_high.pratio, alternative='greater').pvalue
print('Expensive suburbs have less Pupil-teacher ratio by town than cheap ones',
      f'with one sided Mann-Whitney U test p-value: {pval_mann_ptratio}')
```

Expensive suburbs have less Pupil-teacher ratio by town than cheap ones with one sided Mann-Whitney U test p-value: 1.9715109522953002e-15

After testing three most significant predictors I can say that higher cost of suburbs significantly depends on

1. lower status of population
2. number of rooms
3. pupil-teacher ratio

The first parameter is interesting because the population with "low status" simply cannot buy expensive houses. So the correlation is high.

The second parameter is clear - big house = big price.

The third is the dependence on education. The small number of students in one class is characteristic of private schools, which usually educate the children of wealthy parents.

Testing *nox*

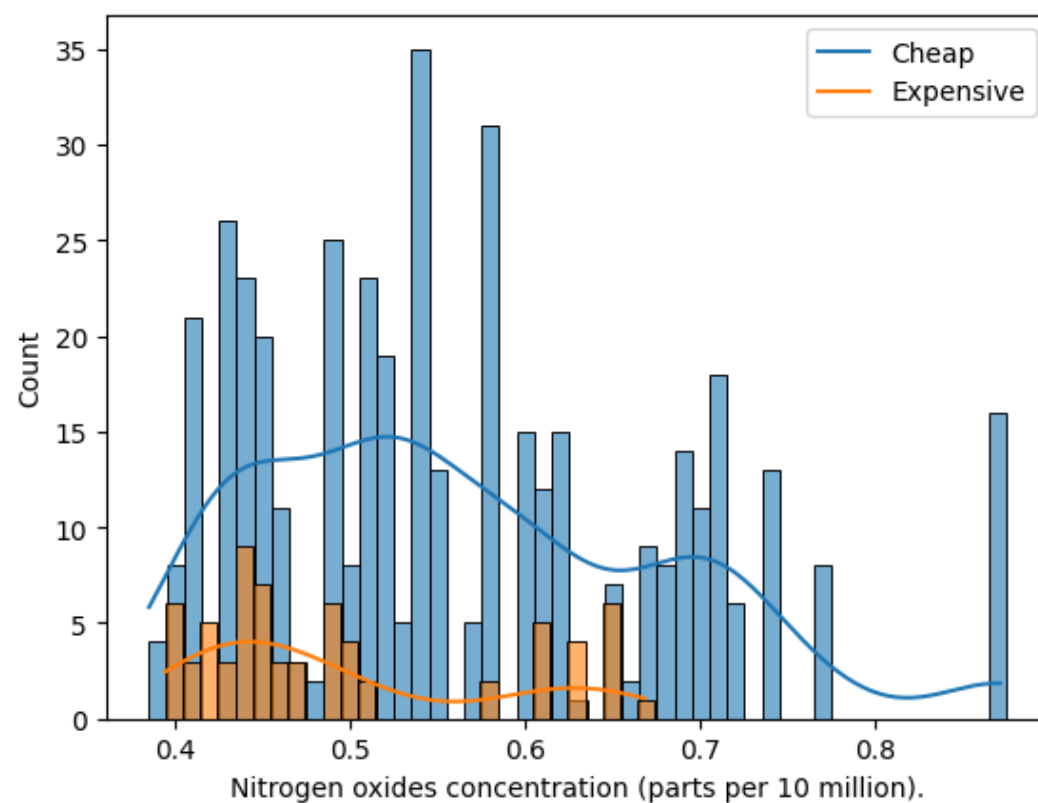
```
In [47]: data_medv_low.nox.describe()
```

```
Out[47]: count    437.000000
mean         0.564281
std          0.117315
min          0.385000
25%          0.464000
50%          0.538000
75%          0.647000
max          0.871000
Name: nox, dtype: float64
```

```
In [48]: data_medv_high.nox.describe()
```

```
Out[48]: count     69.000000
mean         0.493984
std          0.084527
min          0.394000
25%          0.437000
50%          0.458000
75%          0.575000
max          0.668000
Name: nox, dtype: float64
```

```
In [49]: b1 = sns.histplot(data_medv_low.nox, kde=True,
                          alpha=0.6, binwidth=0.01);
b2 = sns.histplot(data_medv_high.nox, kde=True,
                  alpha=0.6, binwidth=0.01);
plt.legend(['Cheap', 'Expensive']);
plt.xlabel('Nitrogen oxides concentration (parts per 10 million).');
# plt.ylim(0,15);
```



```
In [50]: pval_mann_nox = mannwhitneyu(data_medv_low.nox, data_medv_high.nox, alternative='greater').pvalue
print('Expensive suburbs have less Nitrogen oxides concentration (parts per 10 million) than cheap ones',
      f'with one sided Mann-Whitney U test p-value: {pval_mann_nox}')
```

Expensive suburbs have less Nitrogen oxides concentration (parts per 10 million) than cheap ones with one sided Mann-Whitney U test p-value: 3.640067266843077e-07

The same shape of distribution, but different mean value

Looks like it is two groups among all suburbs: with high nitrogen oxides concentration and lowe (two maximum peaks on the distribution)

Testing *crim*

```
In [51]: data_medv_low.crim.describe()
```


```
Out[51]: count    437.000000
mean       4.063478
std        9.147717
min        0.006320
25%        0.092990
50%        0.290900
75%        4.541920
max        88.976200
Name: crim, dtype: float64
```

```
In [52]: data_medv_high.crim.describe()
```

```
Out[52]: count     69.000000
mean       0.763810
std        1.837559
min        0.009060
25%        0.037680
50%        0.086640
75%        0.534120
max        9.232300
Name: crim, dtype: float64
```

```
In [53]: b1 = sns.histplot(data_medv_low.crim, kde=True,
                           alpha=0.6, binwidth=3);
b2 = sns.histplot(data_medv_high.crim, kde=True,
                  alpha=0.6, binwidth=3);
plt.legend(['Cheap', 'Expensive']);
plt.xlabel('Per capita crime rate by town.');
```

```
Out[53]: (0.0, 150.0)
```



```
In [54]: pval_mann_crim = mannwhitneyu(data_medv_low.crim, data_medv_high.crim, alternative='greater').pvalue
print('Expensive suburbs have less per capita crime rate by town than cheap ones',
      f'with one sided Mann-Whitney U test p-value: {pval_mann_crim}')
```

Expensive suburbs have less per capita crime rate by town than cheap ones with one sided Mann-Whitney U test p-value : 6.192291924409378e-07

Cheap districts in some occasions have very high criminality rate (up to 89). Expensive ones have maximum 9 *crim*

Testing *zn*

```
In [55]: data_medv_low.zn.describe()
```


```
Out[55]: count    437.000000
mean         8.599542
std         20.099033
min          0.000000
25%          0.000000
50%          0.000000
75%          0.000000
max         100.000000
Name: zn, dtype: float64
```

```
In [56]: data_medv_high.zn.describe()
```

```
Out[56]: count     69.000000
mean     28.869565
std     33.004529
min      0.000000
25%      0.000000
50%     20.000000
75%     45.000000
max     95.000000
Name: zn, dtype: float64
```

```
In [57]: b1 = sns.histplot(data_medv_low.zn, kde=True,
                           alpha=0.6, binwidth=5);
b2 = sns.histplot(data_medv_high.zn, kde=True,
                  alpha=0.6, binwidth=5);
plt.legend(['Cheap', 'Expensive']);
plt.xlabel('Proportion of residential land zoned for lots over 25,000 sq.ft.');
```

```
Out[57]: (0.0, 60.0)
```





```
In [58]: pval_mann_zn = mannwhitneyu(data_medv_low.zn, data_medv_high.zn, alternative='less').pvalue
print('Expensive suburbs have greater proportion of residential land zoned for lots over 25,000 sq.ft. than cheap on
      f'with one sided Mann-Whitney U test p-value: {pval_mann_zn}')
```

Expensive suburbs have greater proportion of residential land zoned for lots over 25,000 sq.ft. than cheap ones with one sided Mann-Whitney U test p-value: 3.8556018089623235e-12

Scatter

```
In [59]: sns.scatterplot(x=medv, y=scaled_preds.zn);
```



But there are a lot of zero values

Yes expensive houses have less zero *zn*, but it looks like there are simply less expensive houses.

However correlation exists: a lot of such zones => less houses in suburbs and more lands per house => higher cost

Testing indus

```
In [60]: data_medv_low.indus.describe()
```

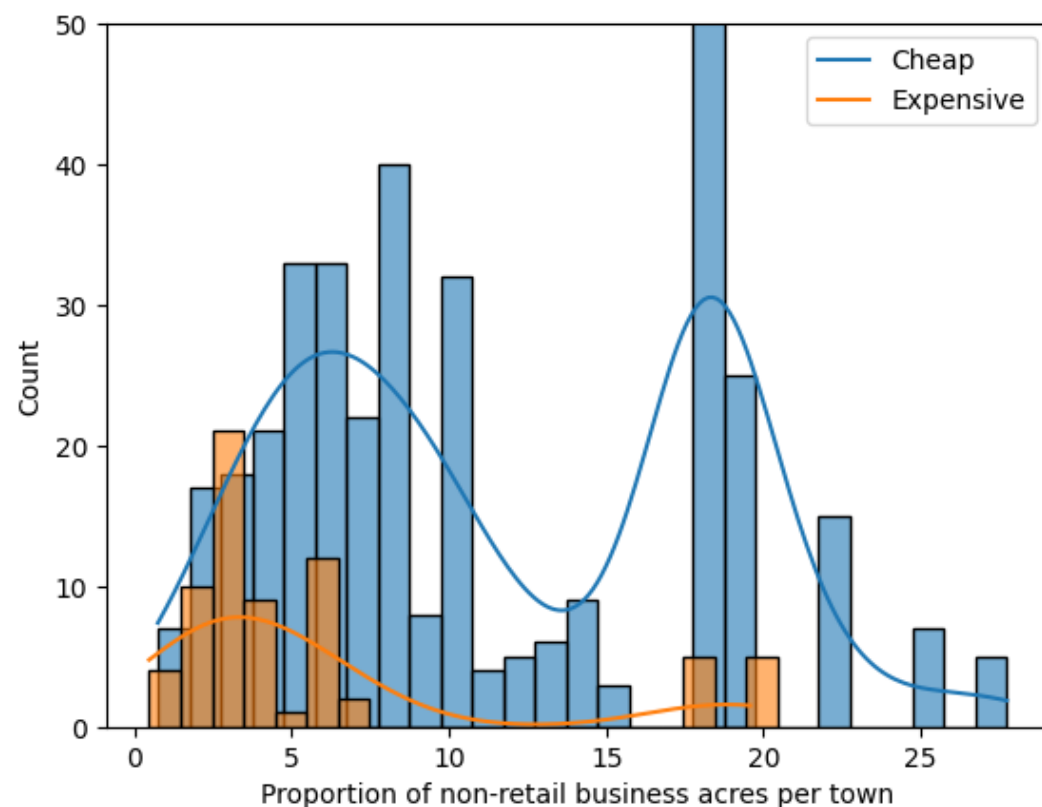
```
Out[60]: count    437.000000
mean       11.975332
std        6.661414
min        0.740000
25%        6.060000
50%       10.010000
75%       18.100000
max       27.740000
Name: indus, dtype: float64
```

```
In [61]: data_medv_high.indus.describe()
```

```
Out[61]: count    69.000000
         mean     5.825942
         std      5.644957
         min      0.460000
         25%      2.460000
         50%      3.440000
         75%      6.200000
         max     19.580000
         Name: indus, dtype: float64
```

```
In [62]: b1 = sns.histplot(data_medv_low.indus, kde=True,
                           alpha=0.6, binwidth=1);
         b2 = sns.histplot(data_medv_high.indus, kde=True,
                           alpha=0.6, binwidth=1);
         plt.legend(['Cheap', 'Expensive']);
         plt.xlabel('Proportion of non-retail business acres per town');
         plt.ylim(0, 50)
```

```
Out[62]: (0.0, 50.0)
```



```
In [63]: pval_mann_indus = mannwhitneyu(data_medv_low.zn, data_medv_high.zn, alternative='less').pvalue
         print('Expensive suburbs have greater proportion of non-retail business acres per town than cheap ones',
               f'with one sided Mann-Whitney U test p-value: {pval_mann_indus}')
```

Expensive suburbs have greater proportion of non-retail business acres per town than cheap ones with one sided Mann-Whitney U test p-value: 3.8556018089623235e-12

The plot looks like expensive houses should have a lower *indus* value, but the Mann-Whitney test shows us the opposite information. I think, that simply number of cheap houses is greater and it affected the results.

Also, it looks like there two groups of suburbs: with high and low values of non-retail business acres.

But some correlation exists!

Testing tax

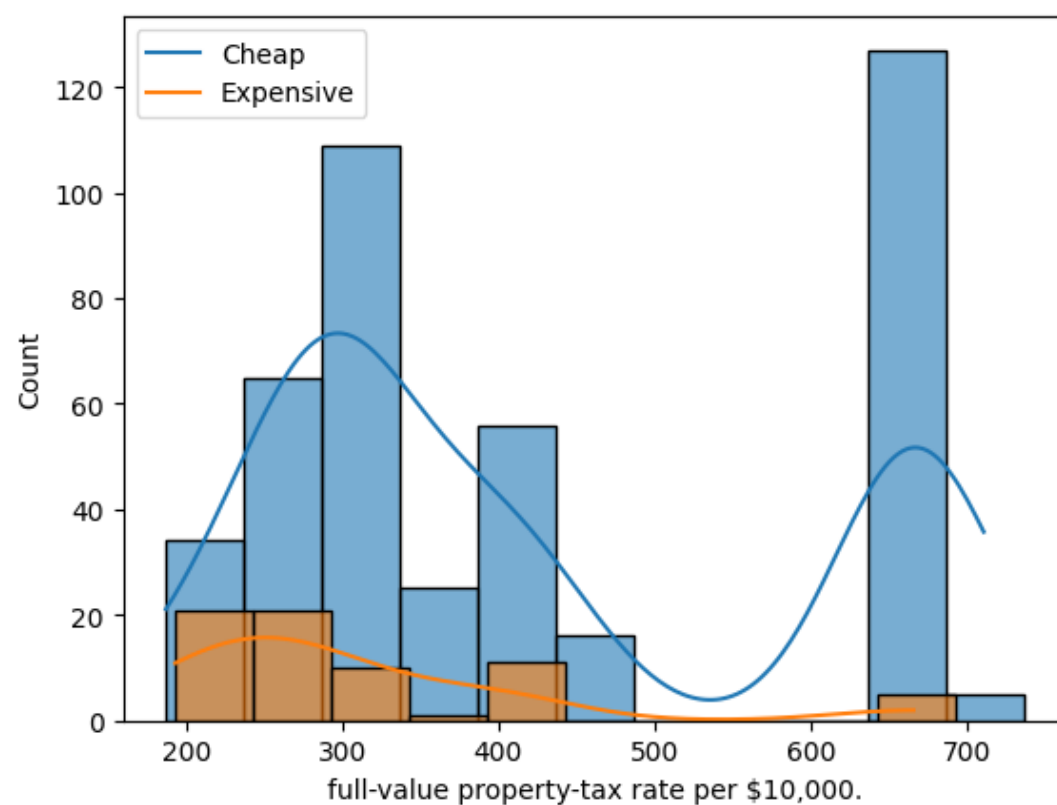
```
In [64]: data_medv_low.tax.describe()
```

```
Out[64]: count    437.000000
         mean    424.109840
         std     169.703456
         min     187.000000
         25%     289.000000
         50%     358.000000
         75%     666.000000
         max     711.000000
         Name: tax, dtype: float64
```

```
In [65]: data_medv_high.tax.describe()
```

```
Out[65]: count     69.000000
         mean    307.710145
         std     120.081725
         min     193.000000
         25%     224.000000
         50%     264.000000
         75%     330.000000
         max     666.000000
         Name: tax, dtype: float64
```

```
In [66]: b1 = sns.histplot(data_medv_low.tax, kde=True,
                        alpha=0.6, binwidth=50);
b2 = sns.histplot(data_medv_high.tax, kde=True,
                  alpha=0.6, binwidth=50);
plt.legend(['Cheap', 'Expensive']);
plt.xlabel('full-value property-tax rate per \$10,000.');
```



The data don't look any different! + there are two peaks that indicate the existence of two groups depending on tax value

CONCLUSION

In my opinion, the most important aspects to consider when choosing an area to build a house are

- Lower status of the population

But dependency there may be inverse relationship - houses in suburbs are expensive and people with lower status cannot buy such houses

- Number of rooms

It is very significant parameter

- Pupil per teacher

Most of rich people want their children to study in private schools with high level of personal education

More complicated parameters:

- Criminality

Criminality in rich suburbs is less on average, but cause of this distribution is because there are several "deviations" in cheap suburbs where criminality rates are extremely high

- Nitrogen oxides concentrationLevel of nitrogen dioxide and non-retail business acres per town

I think this parameters depend on location of various factories and industrial compnies. Such districts are more ecologically friendly. There are no noises, smells and, in fact, people with lower status.

- Proportion of residential land zoned for lots over 25,000 sq.ft

Expensive suburbs have more lands per house, but it looks like there is small correlation. May be there some cheap suburbs that have small population with empty zones without houses.

- Full-value property-tax rate per \$10,000.

This parameter measures cost of public services. But correlation is not very significant. May be it is bacuse some cheap suburbs have big taxes and people live in cheap houses.

Best House ever:

1. With big nuber of rooms
2. With good educational insitutions nearby
3. No industrial complexes nearby
4. Big teritory
5. Good public services in town
6. No criminality in suburb